

## Cost Optimization Overview for Flask SaaS App Deployment

This document describes how the Kubernetes deployment for the Flask SaaS app is cost optimized on DigitalOcean.

### Cost Optimization Strategies

#### 1. Basic Droplets

- The cluster uses basic (smallest recommended) Droplets for worker nodes, minimizing compute costs.
- No additional storage volumes or block storage are attached, reducing storage expenses.

#### 2. Node Count and Autoscaling

- The default node count is set to three, providing high availability while keeping costs low.
- Cluster autoscaling is enabled:
  - Nodes automatically scale up to a maximum of five only when average CPU usage exceeds 70%.
  - This ensures you only pay for extra resources when truly needed, avoiding over-provisioning.

#### 3. Resource Requests and Limits

- Pod resource requests and limits are set appropriately to prevent unnecessary resource allocation.
- This helps the scheduler pack pods efficiently, reducing wasted capacity.

#### 4. Load Balancer Usage

- Only one LoadBalancer service is provisioned, minimizing cloud networking costs.

#### 5. Stateless Application

- The Flask app is stateless, so no persistent storage is required.
- This allows for easy scaling and avoids costs associated with stateful sets or database hosting.

#### 6. Monitoring and Alerts

- Resource usage is monitored, and alerts are set for abnormal spikes.
- This enables proactive scaling and cost management.

#### 7. Image Optimization

- The Docker image uses a slim Python base, reducing image size and speeding up deployments, which can lower bandwidth and storage costs.

## 8. Idle Resource Management

- Unused resources (nodes, services) are regularly reviewed and decommissioned to avoid unnecessary charges.

### Summary

This deployment is designed to balance reliability and cost:

- Three nodes by default for redundancy.
- Autoscaling up to five nodes only when needed.
- No extra storage or unnecessary services.
- Efficient resource allocation and monitoring.

These practices help ensure the Flask SaaS app runs reliably while keeping cloud costs predictable